Deep learning architectures

Notes

The same ingredients

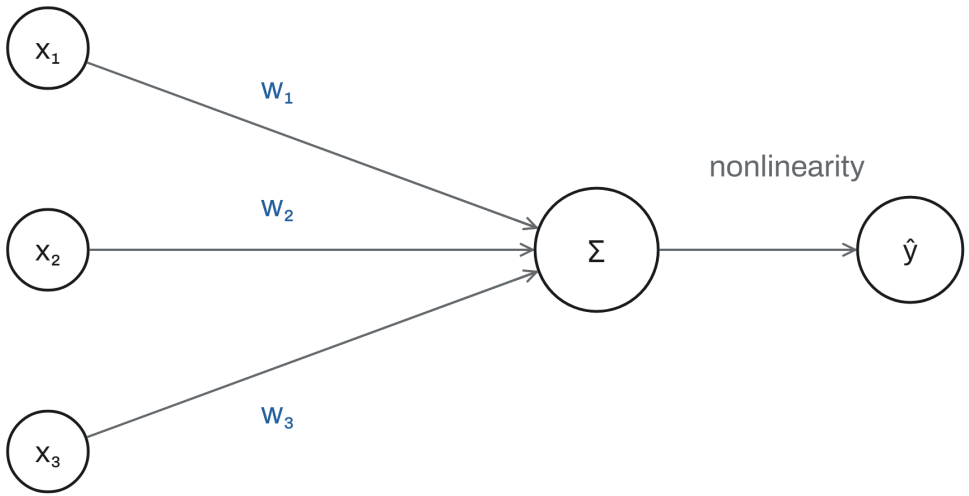

Inputs contribute through weights.

A nonlinearity turns the sum into an output.

Training changes the weights to reduce prediction error.

Notes

Architecture changes the information flow

Layers

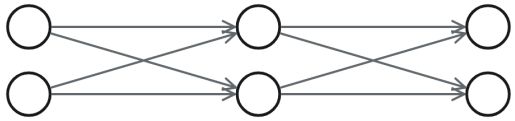

Layers combine features.

Recurrence

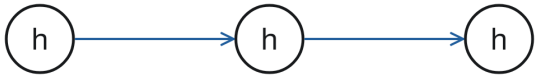

Recurrence carries information through time.

Convolution

shared filter

Convolution reuses a local pattern.

Attention

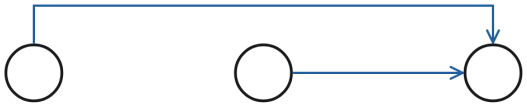

weighted context

Attention mixes information from context.

Notes

A hidden layer creates a representation

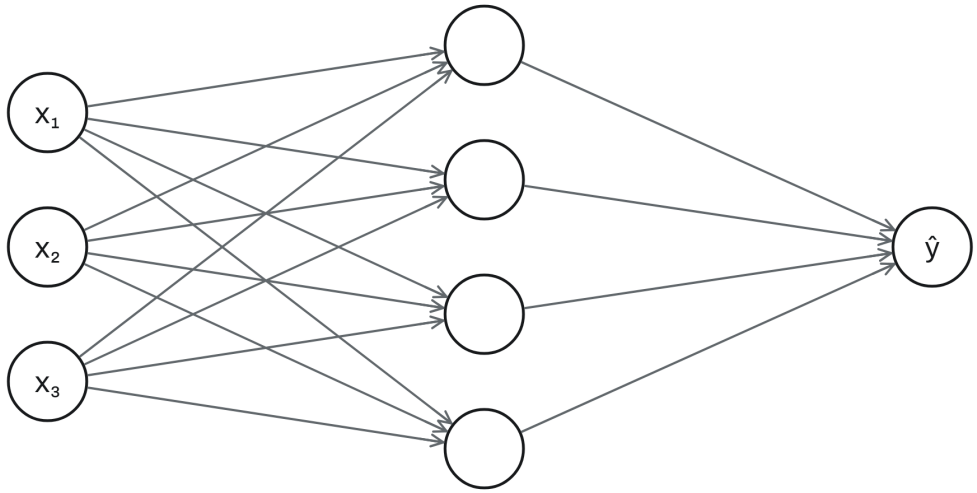

A hidden layer transforms the inputs.

The output uses that transformed pattern.

If we freeze the hidden layer, what can training still change?

The output weights and bias can change. The hidden representation stays fixed.

Notes

A multilayer perceptron (MLP)

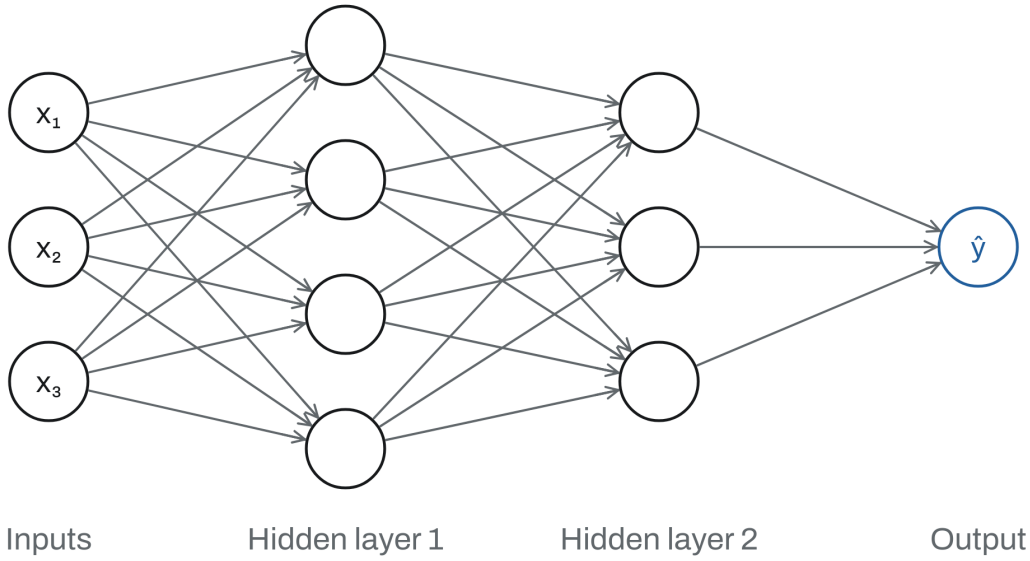

The next layer combines features from the previous layer.

Depth means repeated transformation.

Notes

XOR: one straight boundary is not enough

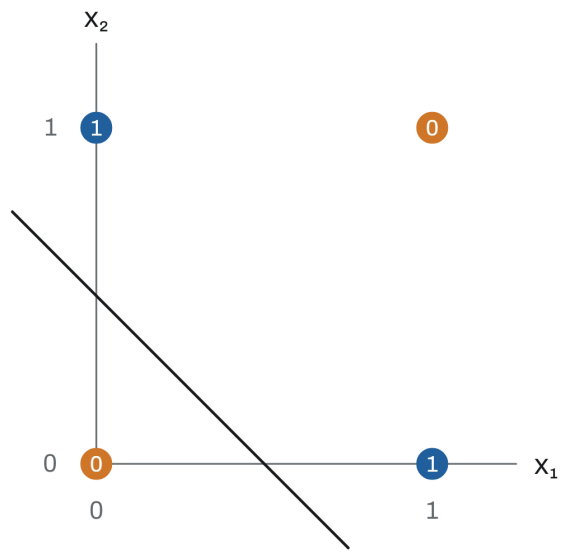

0: same inputs

1: different inputs

Output 1 when exactly one input is 1.

This straight boundary gets three of the four cases right.

Moving the line cannot separate the two diagonals.

Notes

Two hidden units ask simpler questions

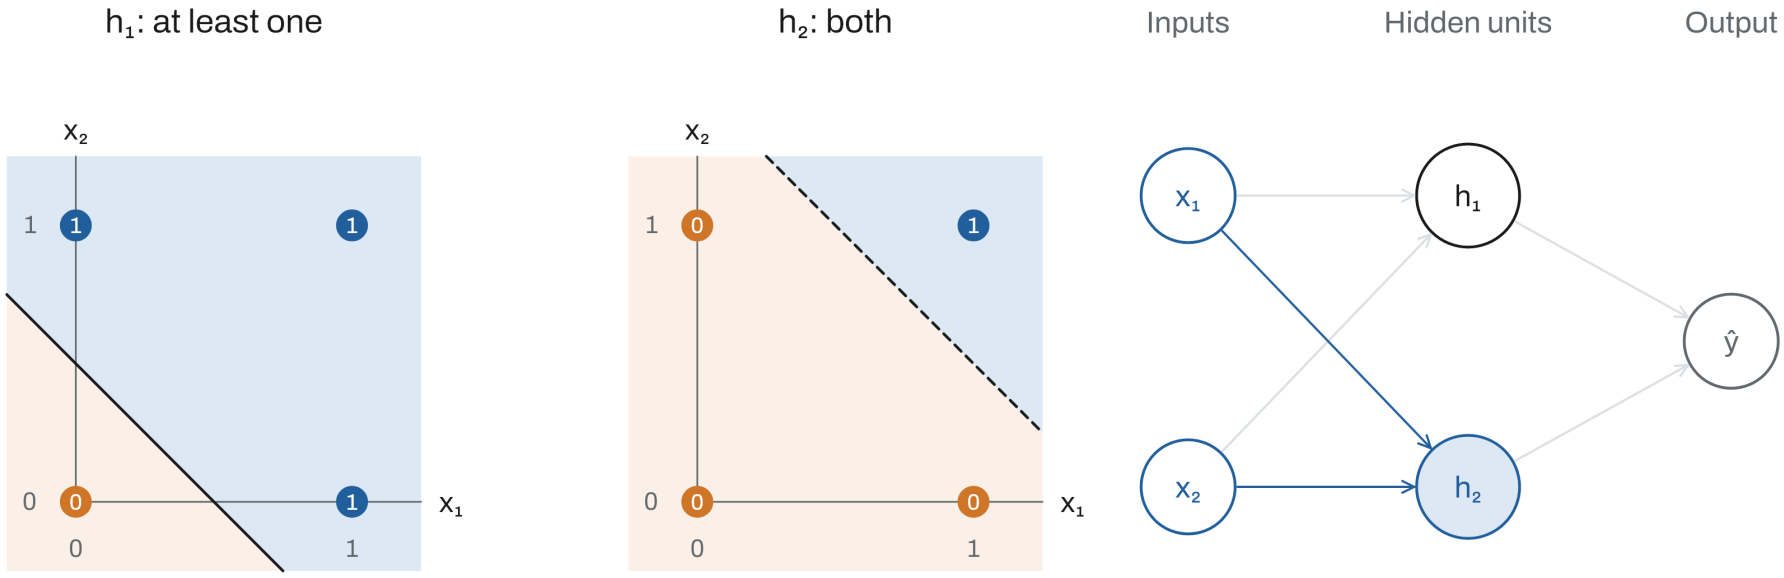

Step-function hidden units · colors show each unit's output

Notes

Combine the hidden decisions to solve XOR

Output decision region

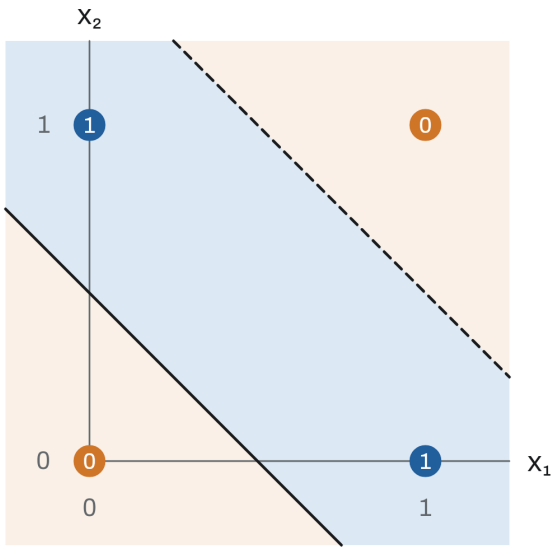

0: same inputs

1: different inputs

Exactly one input → output 1.

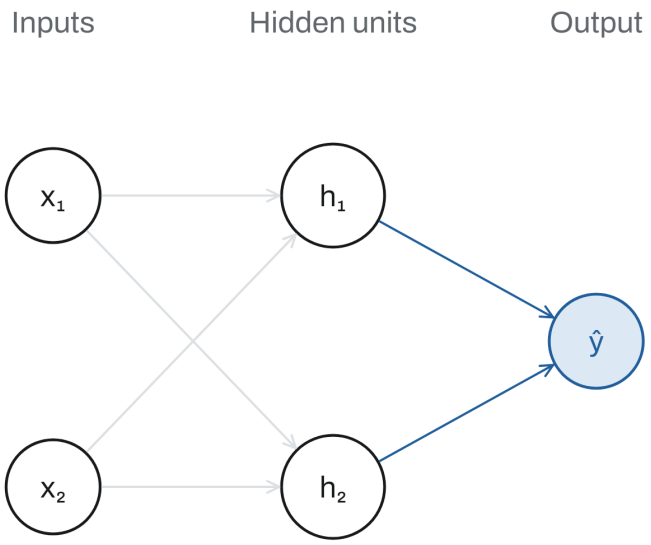

Notes

Learn this decision boundary with two hidden units

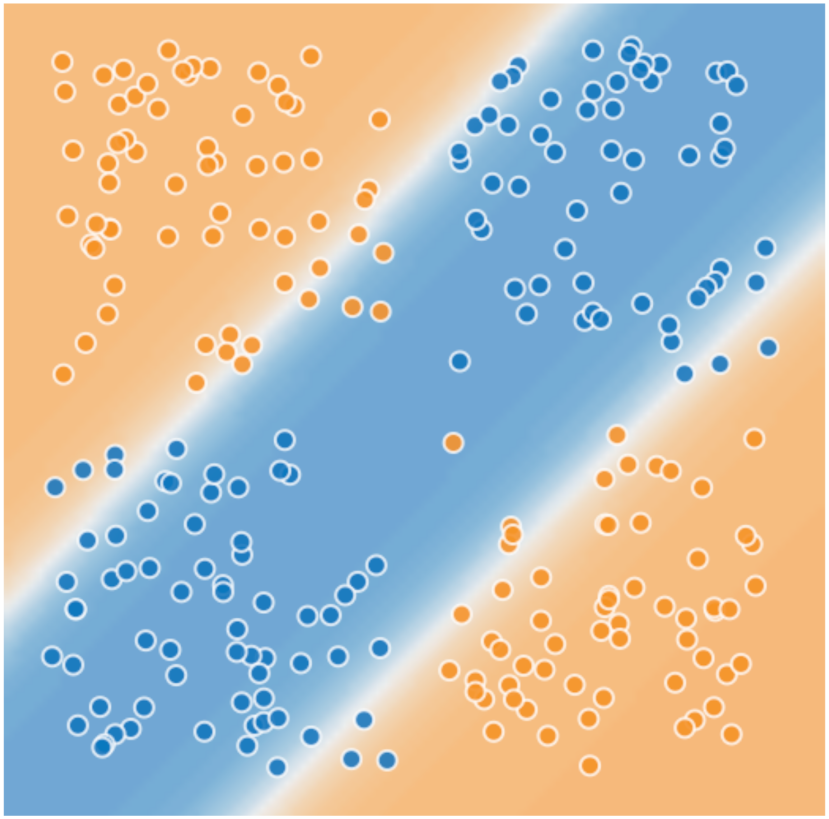

Neural Network Playground

Notes

Semantic memory

robin

has feathers

can fly

canary

has feathers

can fly

rose

has petals

can grow

daisy

has petals

can grow

Notes

Concepts have a similarity structure

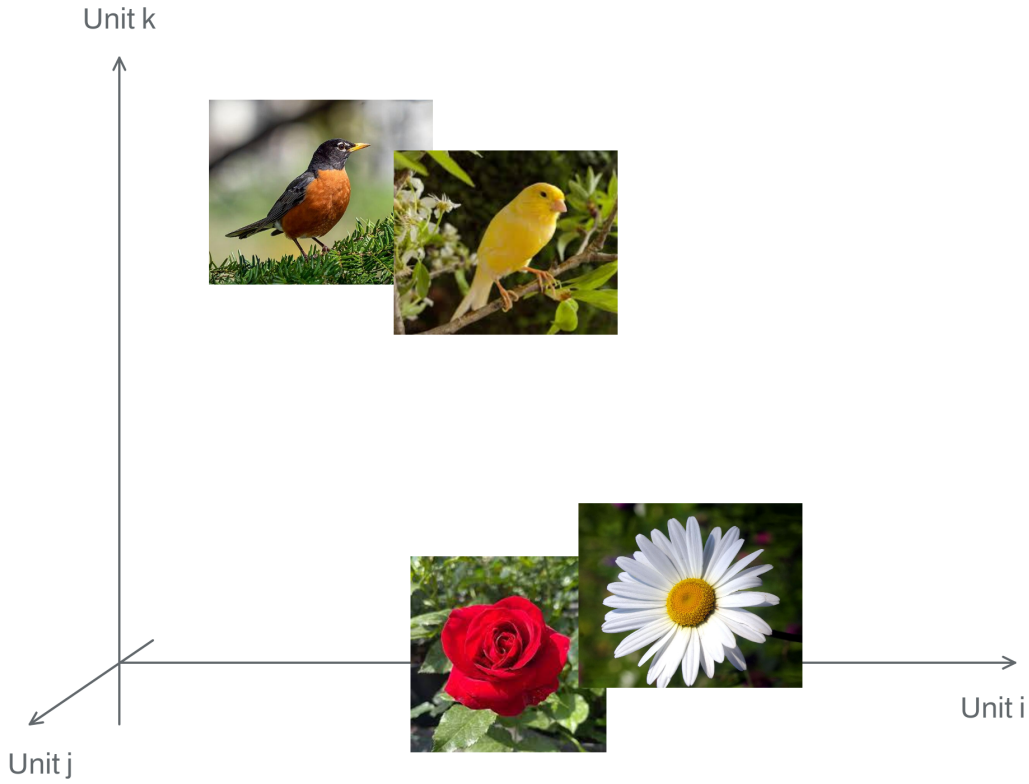

Similar concepts evoke similar neural patterns.

Notes

Concepts have a similarity structure

Similar concepts evoke similar neural patterns.

New Songbird Species Discovered in Bolivia

Jan 5, 2026 by Natali Anderson

« Previous | Next »

Published in
Biology
Featured

Tagged as
Beni greenlet
Bird
Bolivia
DNA
Gene

Deep in the seasonally flooded savannas of Bolivia, a small, olive-green songbird has been quietly eluding science. Now, after six decades of confusion and misidentification, ornithologists have confirmed that the bird is not just a regional variant of a known species in the greenlet genus *Hylophilus* but an entirely new one. The discovery adds a new endemic species to South America's already rich avian diversity and underscores how much remains unknown even in well-studied bird families.

Notes

Broad and fine semantic distinctions

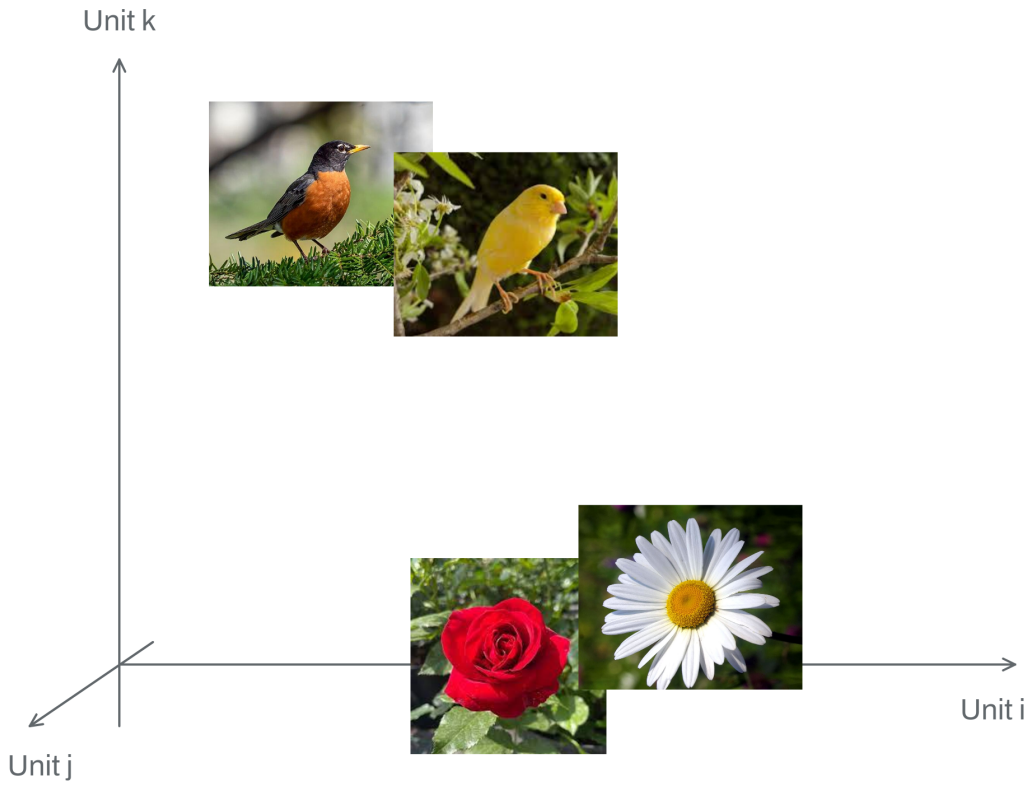

Development: broad categories can emerge before finer distinctions.

Mandler & McDonough (1993)

Semantic dementia: distinctive knowledge can be more vulnerable.

Rogers et al. (2004)

What learning process could produce this structure?

Notes

Semantic knowledge as a prediction task

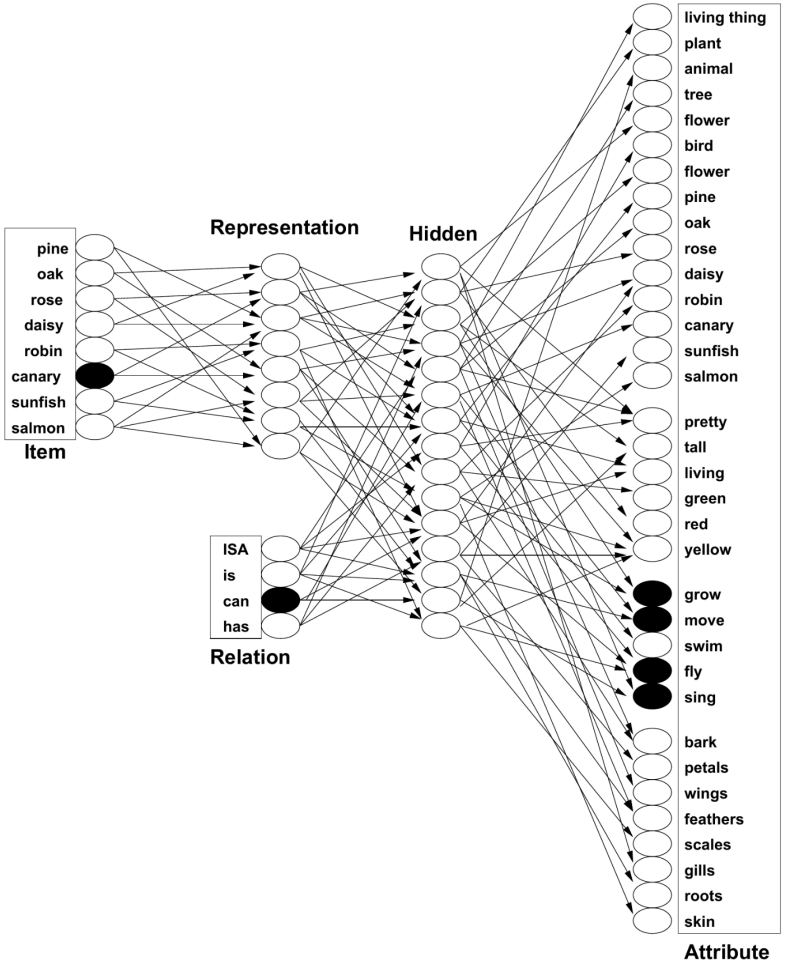

Eight item inputs and four relation inputs.

One item and one relation are selected for each query.

“Canary” + “can” → grow, move, fly, sing.

Rogers & McClelland (2004), Semantic Cognition · Chs. 2–3 · Figure 2.2

Notes

Read a distributed representation

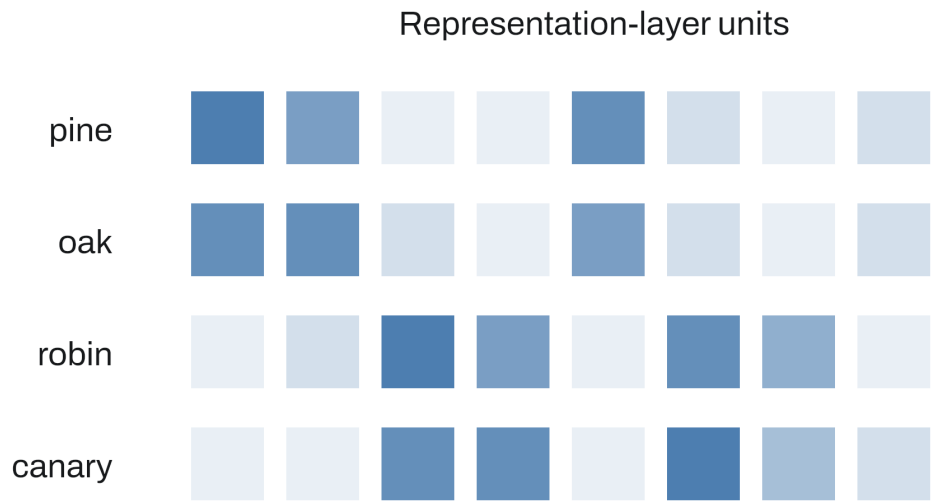

Illustrative values · darker = more activity

Each column is a Representation-layer unit.

Each row shows the unit activations for one item.

Similar patterns can support similar predictions.

Does a similar pattern require one “plant neuron”?

No. Similarity can be distributed across several units.

Notes

Learning separates semantic structure

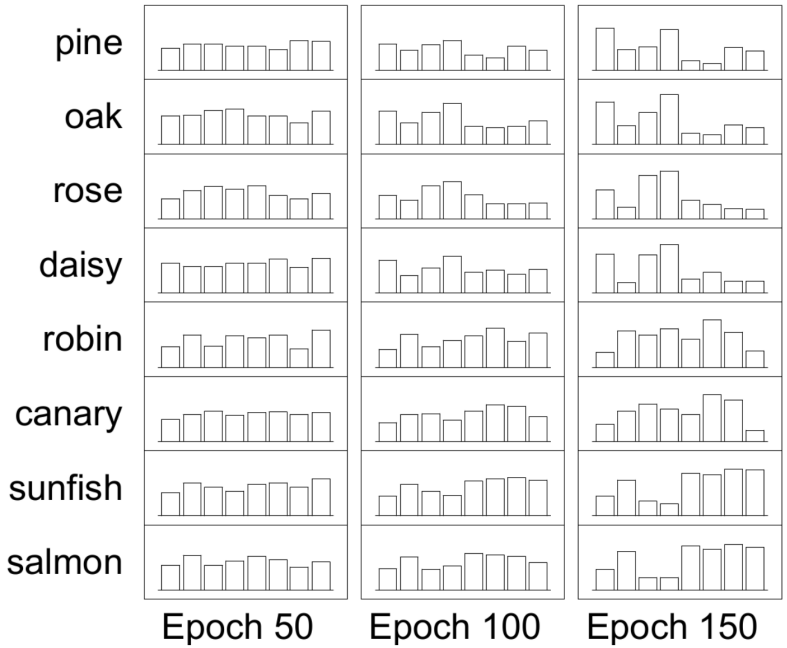

Representation layer: eight learned units.

Each bar is one unit’s activation for the item in its row.

Columns: training epochs 50, 100 and 150.

Related items develop similar activation patterns.

The task and training experience shape this semantic structure.

Rogers & McClelland (2004), Semantic Cognition · Chs. 2–3

Notes

Learning differentiates concepts gradually

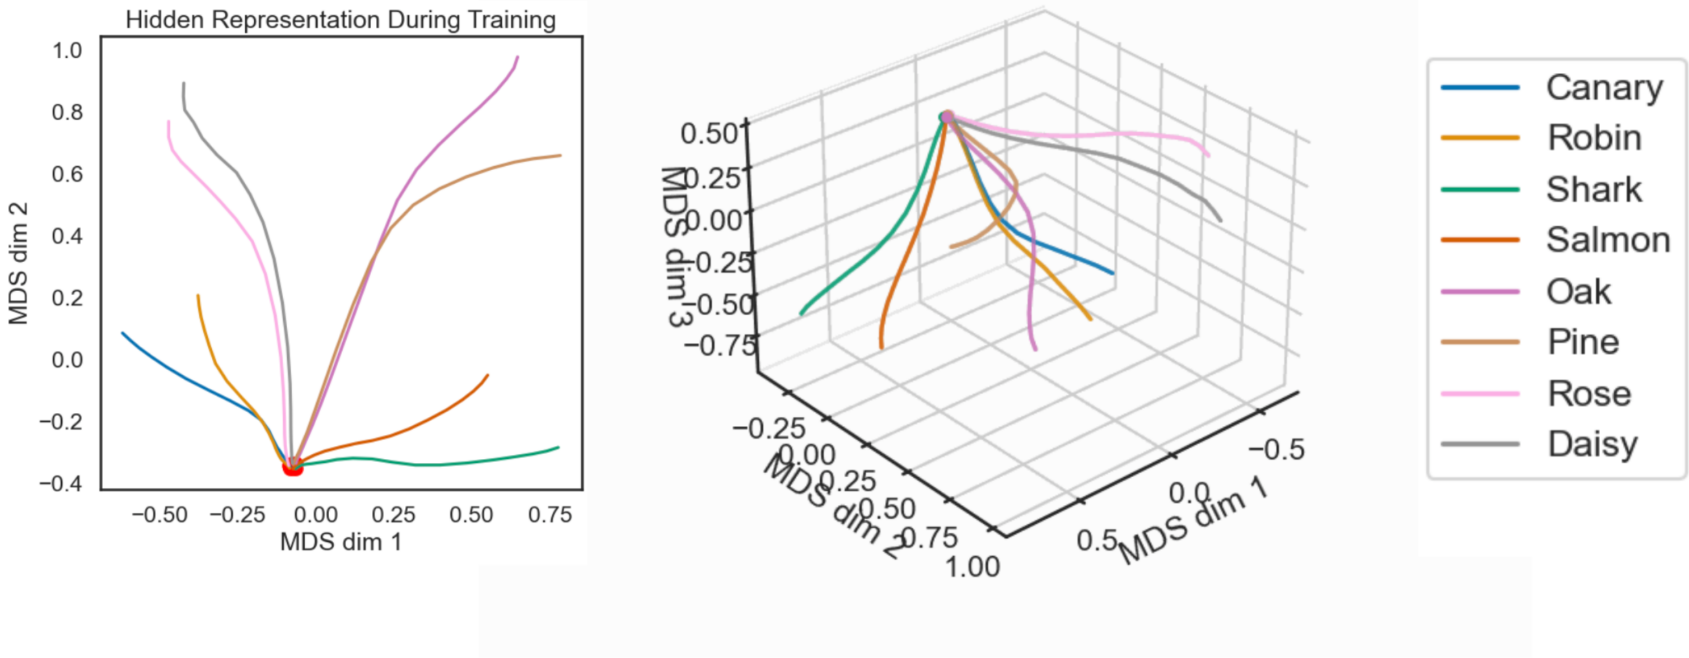

Learning separates broad groups and then finer distinctions

Notes

The same word can need different context

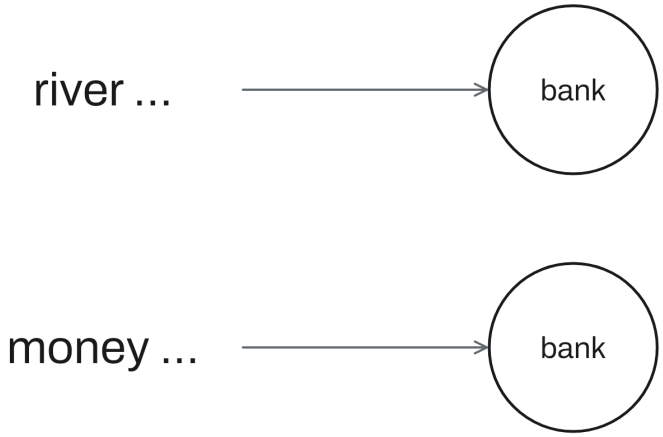

“She walked along the river bank.”

“She deposited money at the bank.”

Context helps distinguish the meanings.

Notes

Recurrent neural network (RNN)

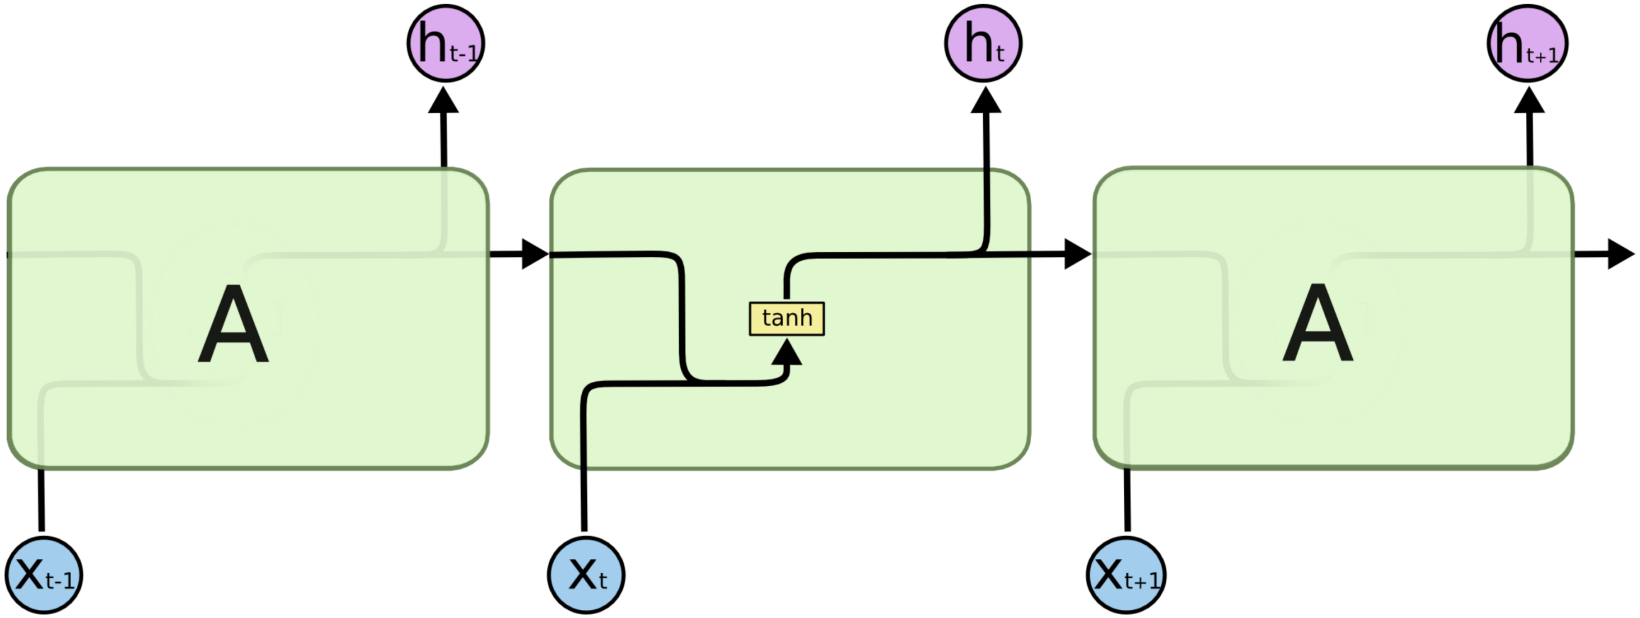

The new hidden state uses the input and the old hidden state.

The state carries information from earlier inputs.

Christopher Olah (2015) · Understanding LSTM Networks

Notes

Unroll the network through time

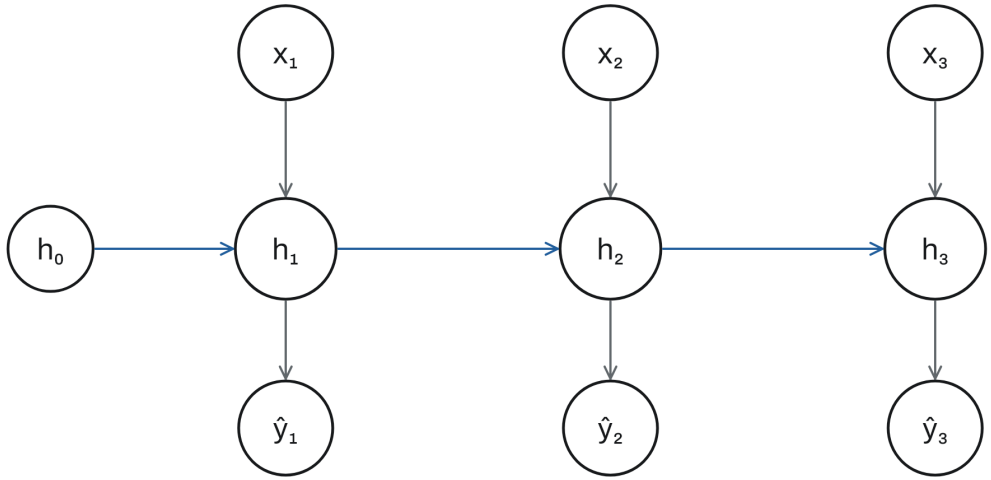

Shared weights across time

One recurrent network, shown at several moments.

The same weights are reused at every step.

Christopher Olah (2015) · Understanding LSTM Networks

Notes

Watch a context update

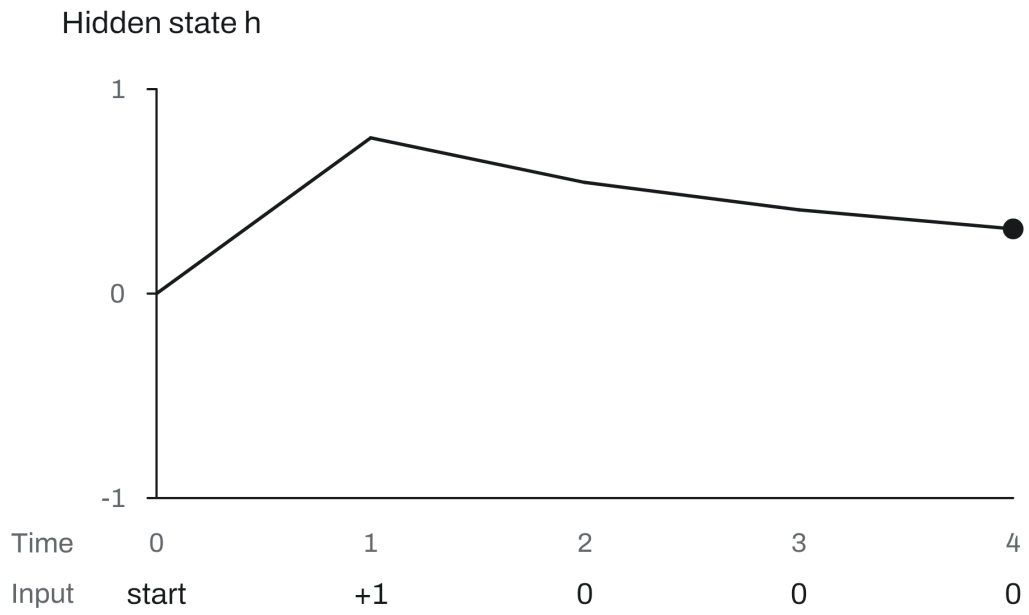

$h = \tanh(\text{input} + 0.8 \times \text{previous } h)$

Compare a first input of +1 with a first input of −1. All later inputs are 0.

First input: +1 · Step 4

$h = 0.316$

$\tanh(0 + 0.8 \times 0.409) = 0.316$

After input 0, is the state the same if we started with +1 or −1?

No. The first input still affects the hidden state.

Notes

Finding word boundaries without pauses

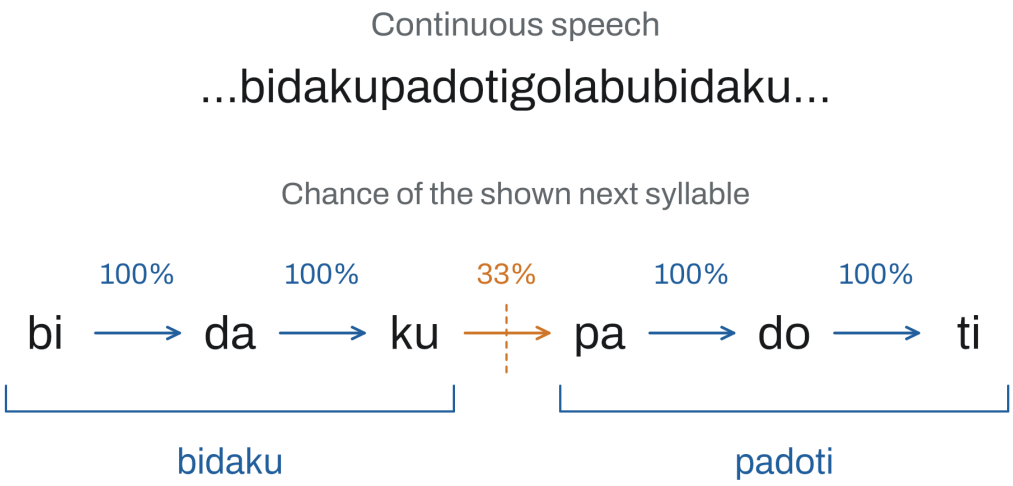

A drop in predictability can signal a boundary

8-month-old infants heard two minutes of continuous artificial speech.

Syllables within a word were more predictable than syllables across words.

At test, infants distinguished words from sequences crossing a boundary.

Saffran, Aslin & Newport (1996) · Statistical Learning by 8-Month-Old Infants

Notes

Elman: predict the next word

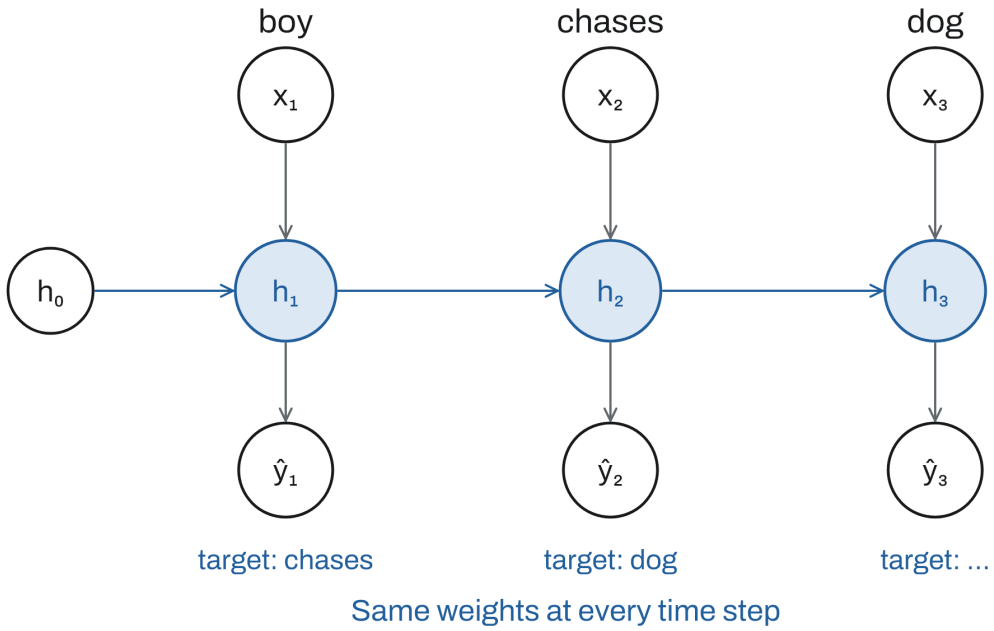

Input: one word at a time.

Training target: the following word.

Hidden states can reflect distributional structure.

Elman (1990), Finding Structure in Time · Cognitive Science

Notes

Structure appears in hidden states

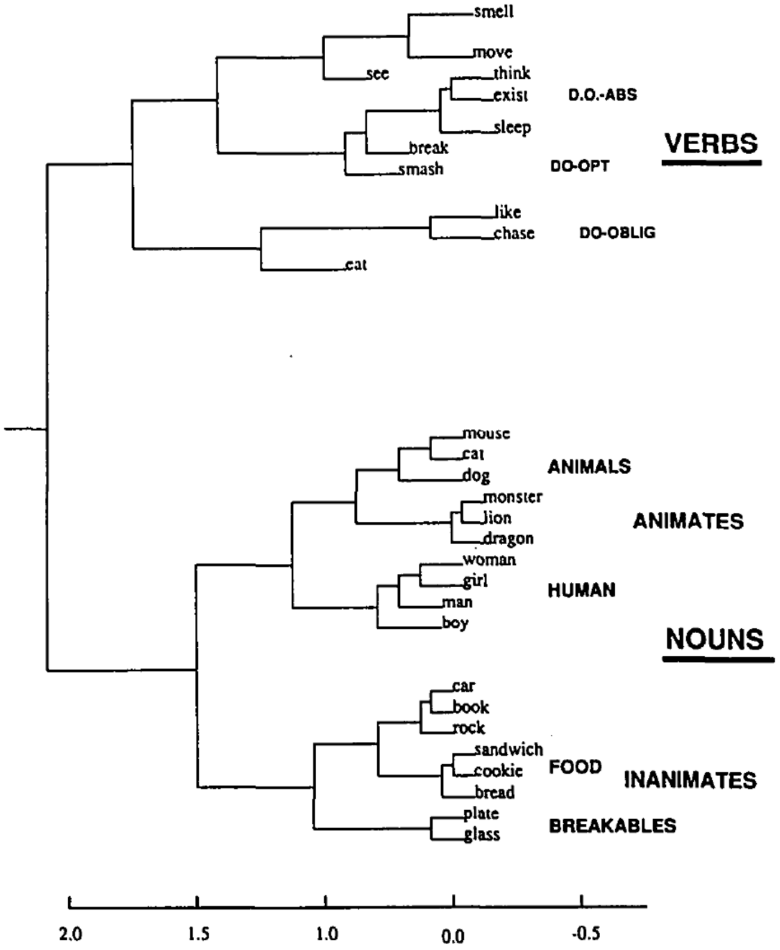

Words clustered by the RNN's hidden-layer activity.

Each word is represented by an average pattern across 150 hidden units.

The average combines the word's occurrences in different sentence contexts.

Words joining nearer the labels have more similar patterns.

Elman (1990), Finding Structure in Time · Cognitive Science

Notes

Human recall favors nearby studied items

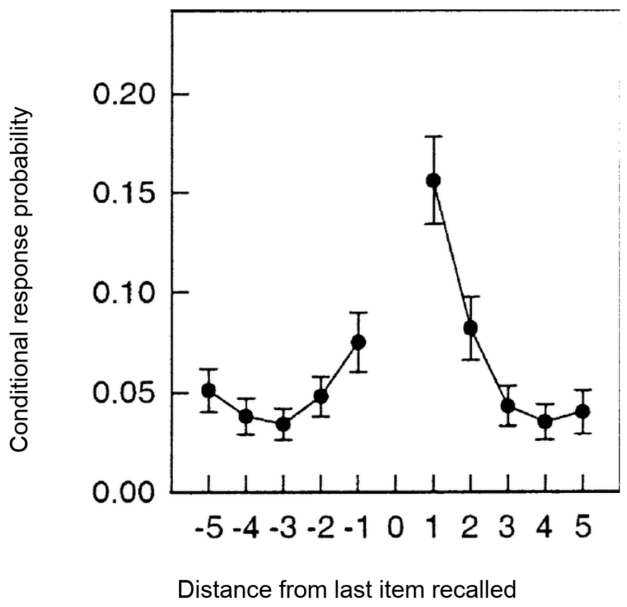

People study a list, then freely recall its items.

Horizontal axis: positions apart in the studied list.

+1 means the item studied immediately after the one just recalled.

Nearby items are more likely to be recalled next.

Howard & Kahana (2002), Temporal Context Model · J. Mathematical Psychology

Notes

Temporal context explains nearby recall

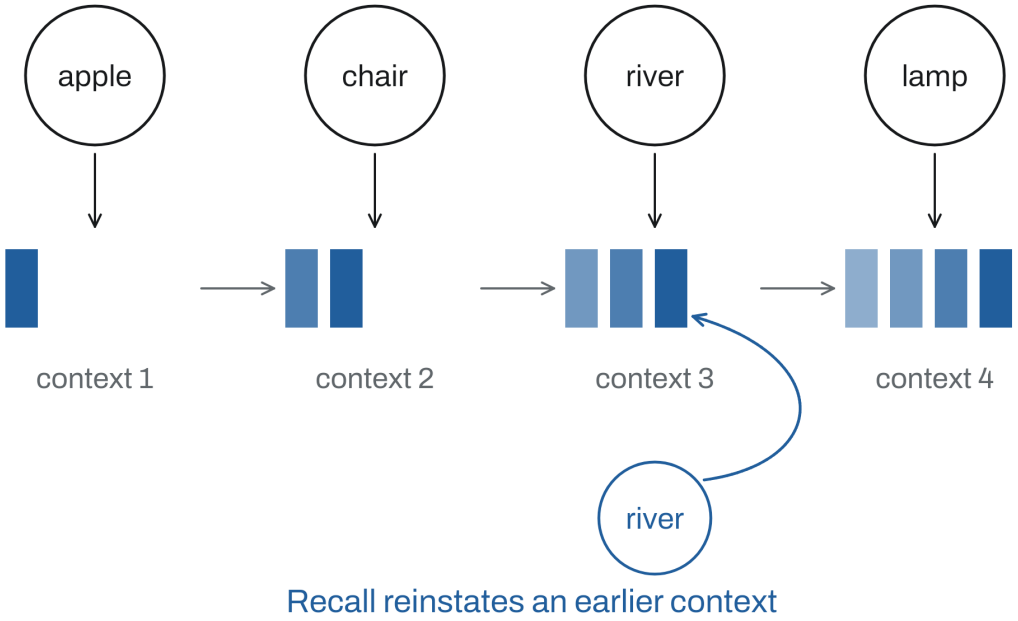

Each bar group is a context vector; darker bars mean stronger components.

Nearby studied items acquire overlapping contexts.

Recall reinstates context, which cues nearby items.

If recalling “river” brings back its context, which items gain a useful cue?

Items studied near “river” share more of that context.

Howard & Kahana (2002), Temporal Context Model · J. Mathematical Psychology

Notes

Keeping a cue across a delay

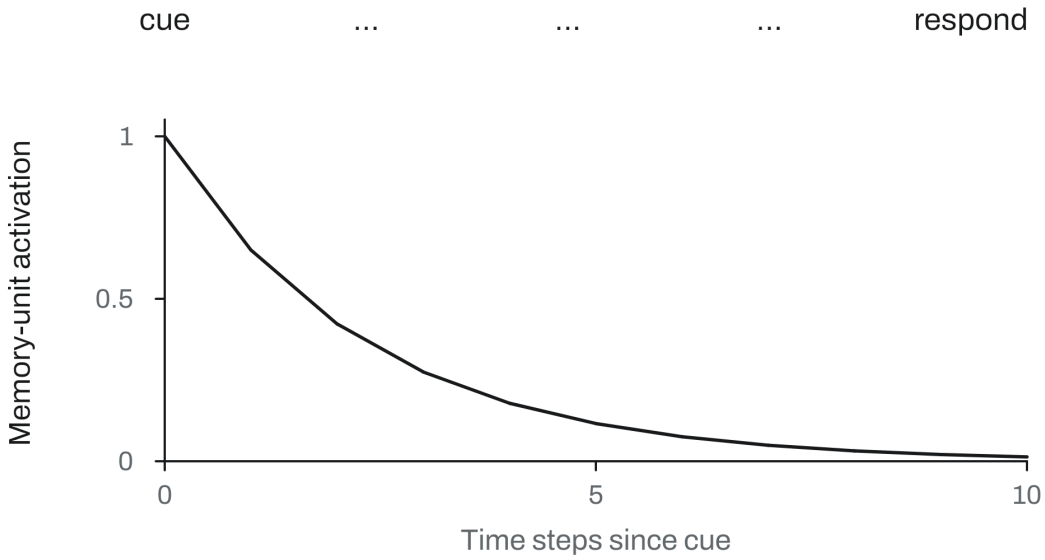

Toy memory unit: start at 1; retain 65% each step

The cue sets one memory unit's activation to 1.

With no new input, each step retains 65% of its previous activation.

The cue's signal becomes weaker across the delay.

Notes

A gate is a learned amount

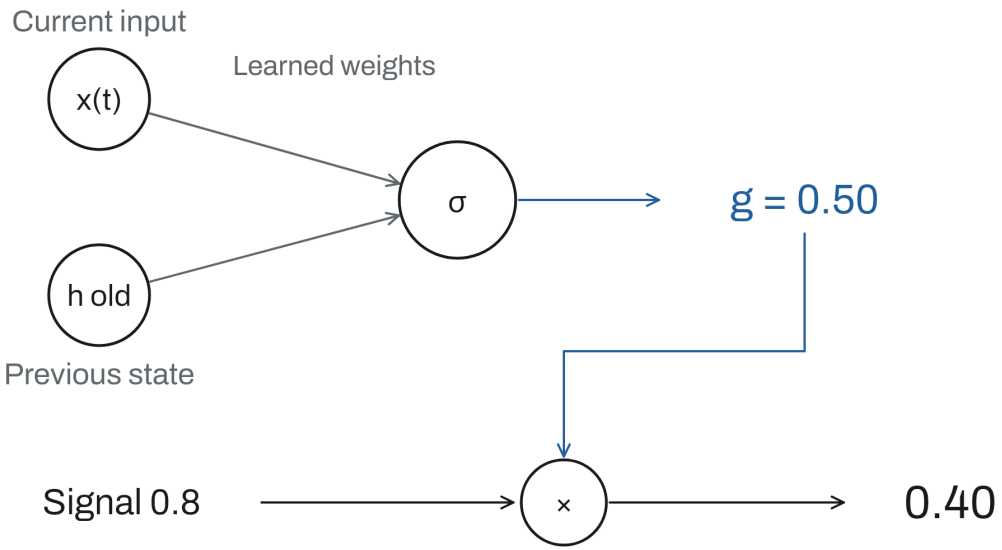

Current input and previous state determine the gate.

A sigmoid turns their weighted sum into an amount from 0 to 1.

Multiplication scales how much a signal contributes.

Gate g: 0.50

Hand-set gate for demonstration

Notes

Long short-term memory (LSTM)

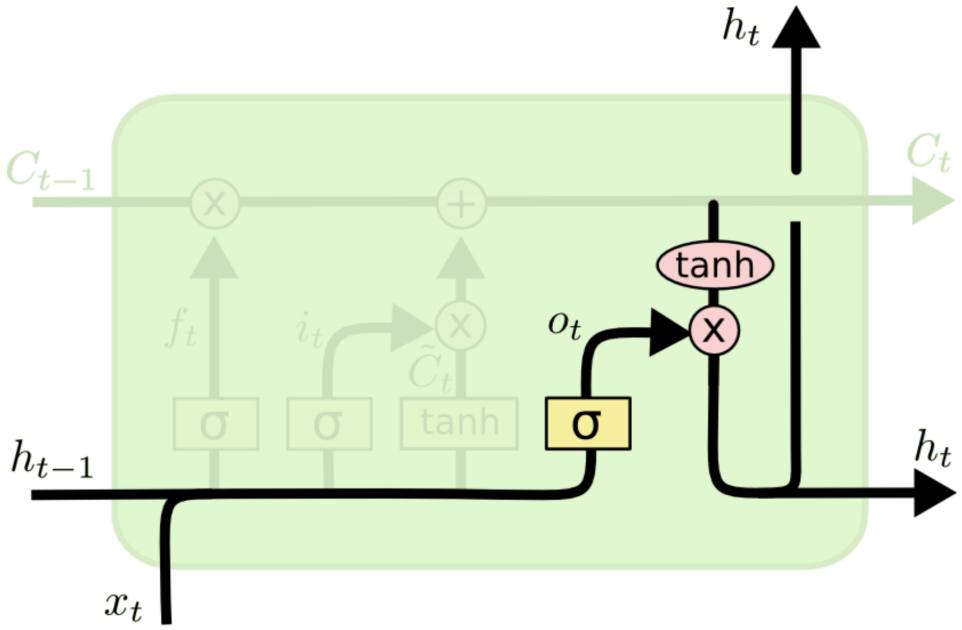

The cell state carries information forward.

The forget gate controls what to retain.

The input gate controls what to write.

Combine retained information with the new contribution.

The output gate controls the visible hidden state.

Can a stored cell value differ from the visible hidden output?

Yes. The output gate controls how much of tanh(cell) is exposed.

Christopher Olah (2015) · Diagram; Hochreiter & Schmidhuber (1997); Gers et al. (2000)

Notes

Preserve a value, then replace it

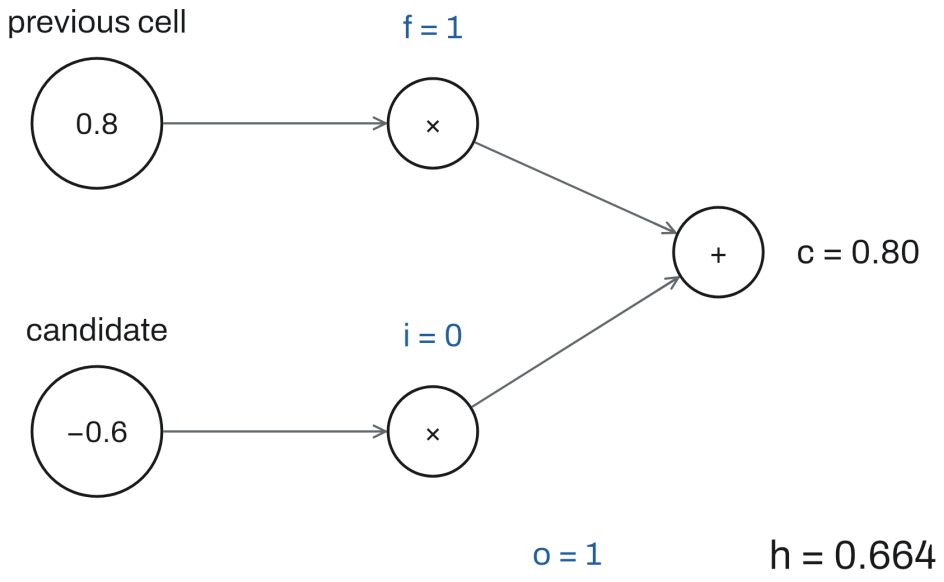

$c = f \times \text{previous } c + i \times \text{candidate}$

$h = o \times \tanh(c)$

c = 0.800

h = 0.664

Notes

Gated recurrent unit (GRU)

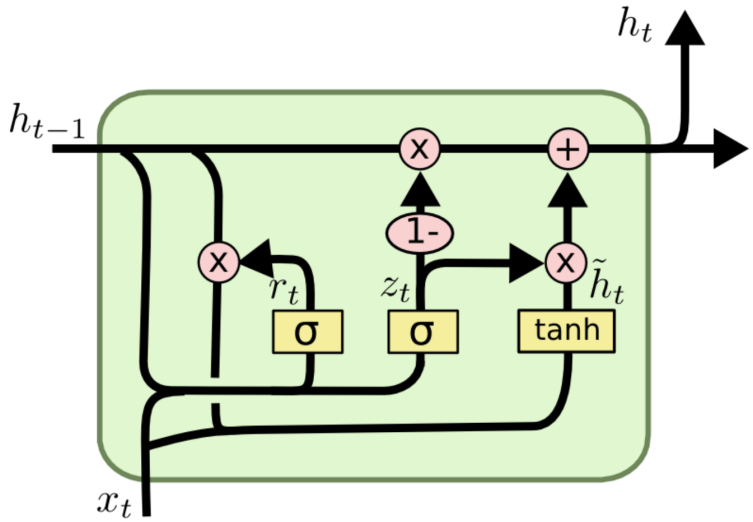

Old state 0.80 → new state 0.42

Candidate 0.04 · update gate writes 50%

One recurrent state, with no separate cell state.

The update gate controls the old/new blend.

Reset changes how the old state forms the candidate.

Update z (write): 0.50

Reset r: 1.00

Set writing to 0. Does changing reset change the stored state?

No. With writing at 0, the old state is retained regardless of the candidate.

Christopher Olah (2015) · Diagram; Cho et al. (2014)

Notes

Same bird, different location

Notes

A visual pattern can move

Here

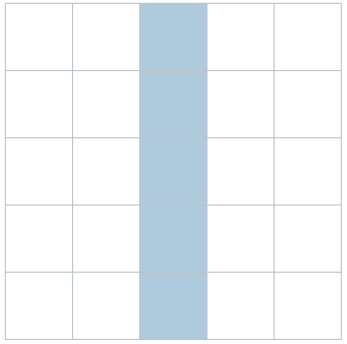

And here

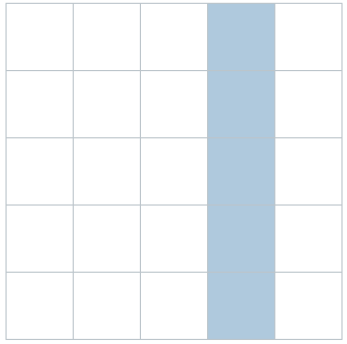

The local pattern is the same.

The same edge can appear in different places.

A dense network gives each location separate weights.

A convolution reuses a local detector.

Notes

One patch produces one number

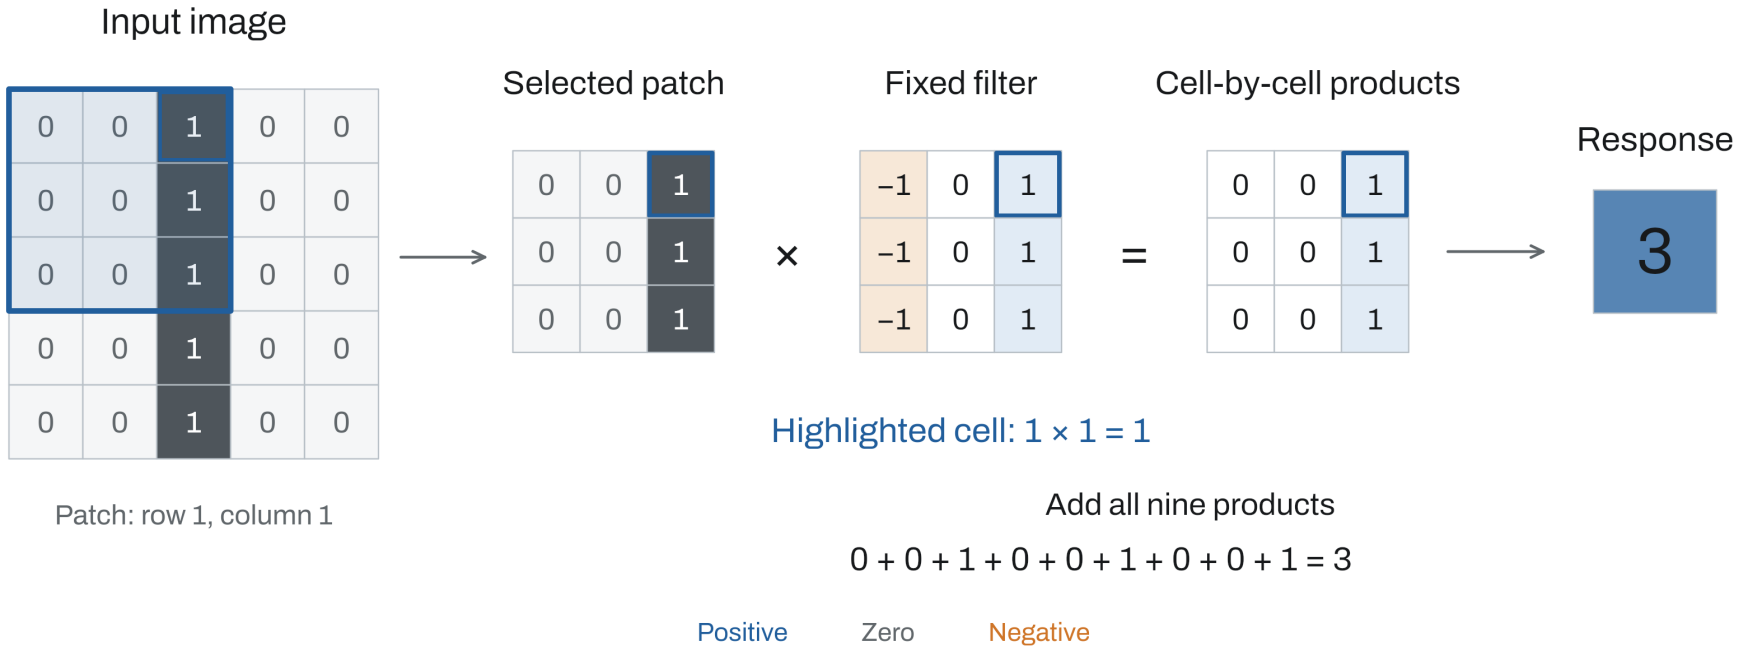

Notes

Build a feature map one position at a time

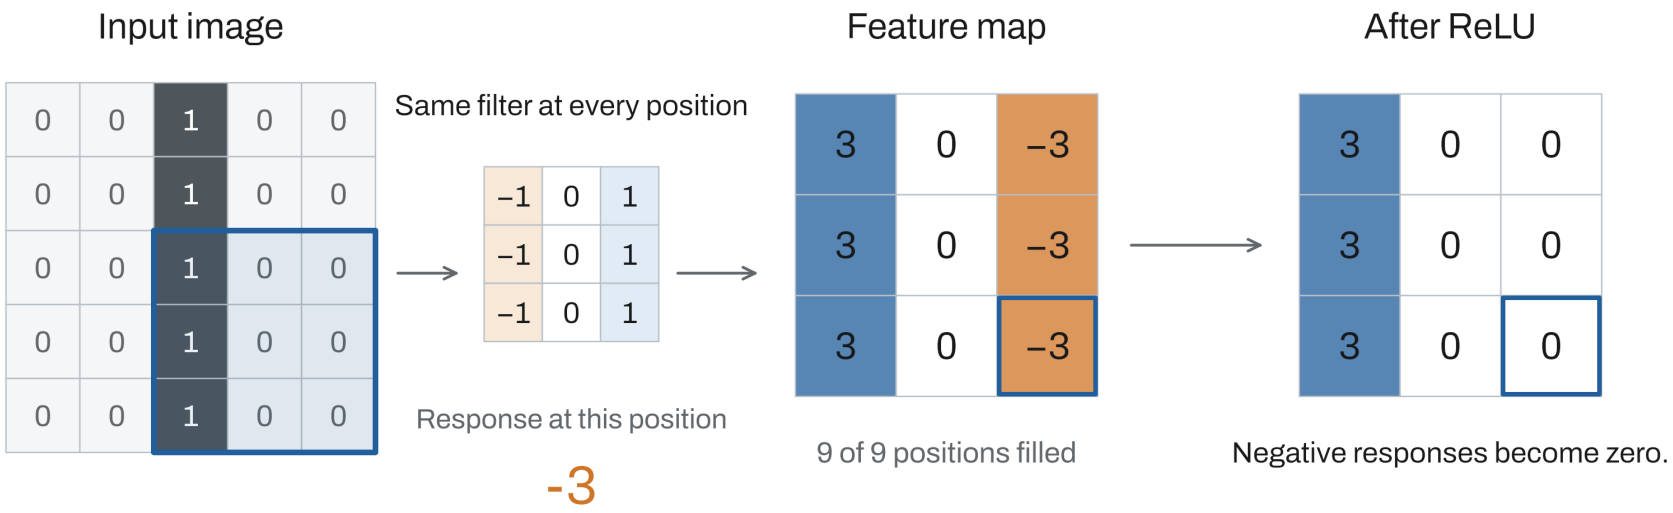

Each move fills one corresponding output cell.

Positive Zero Negative

Notes

Two filters, two different feature maps

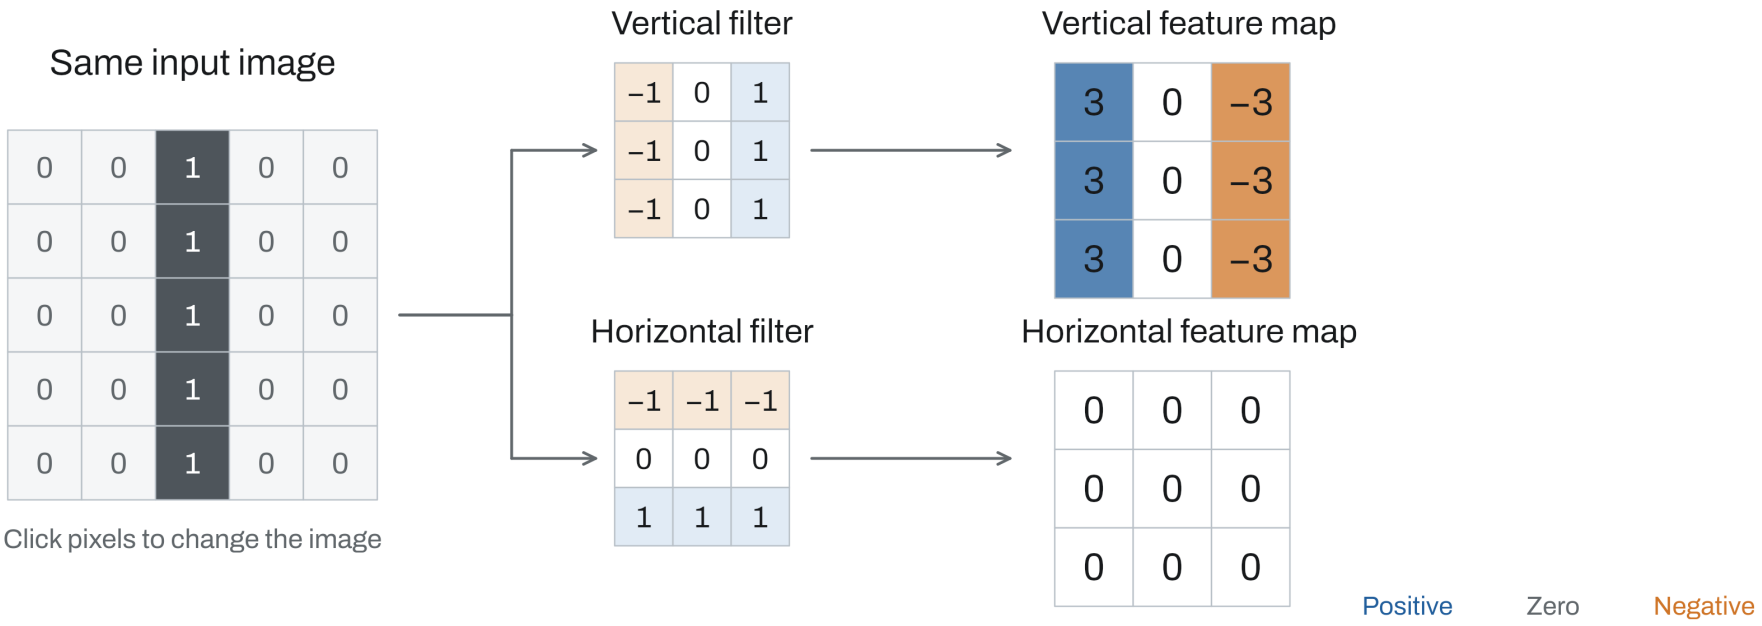

Notes

AlexNet: a convolutional neural network

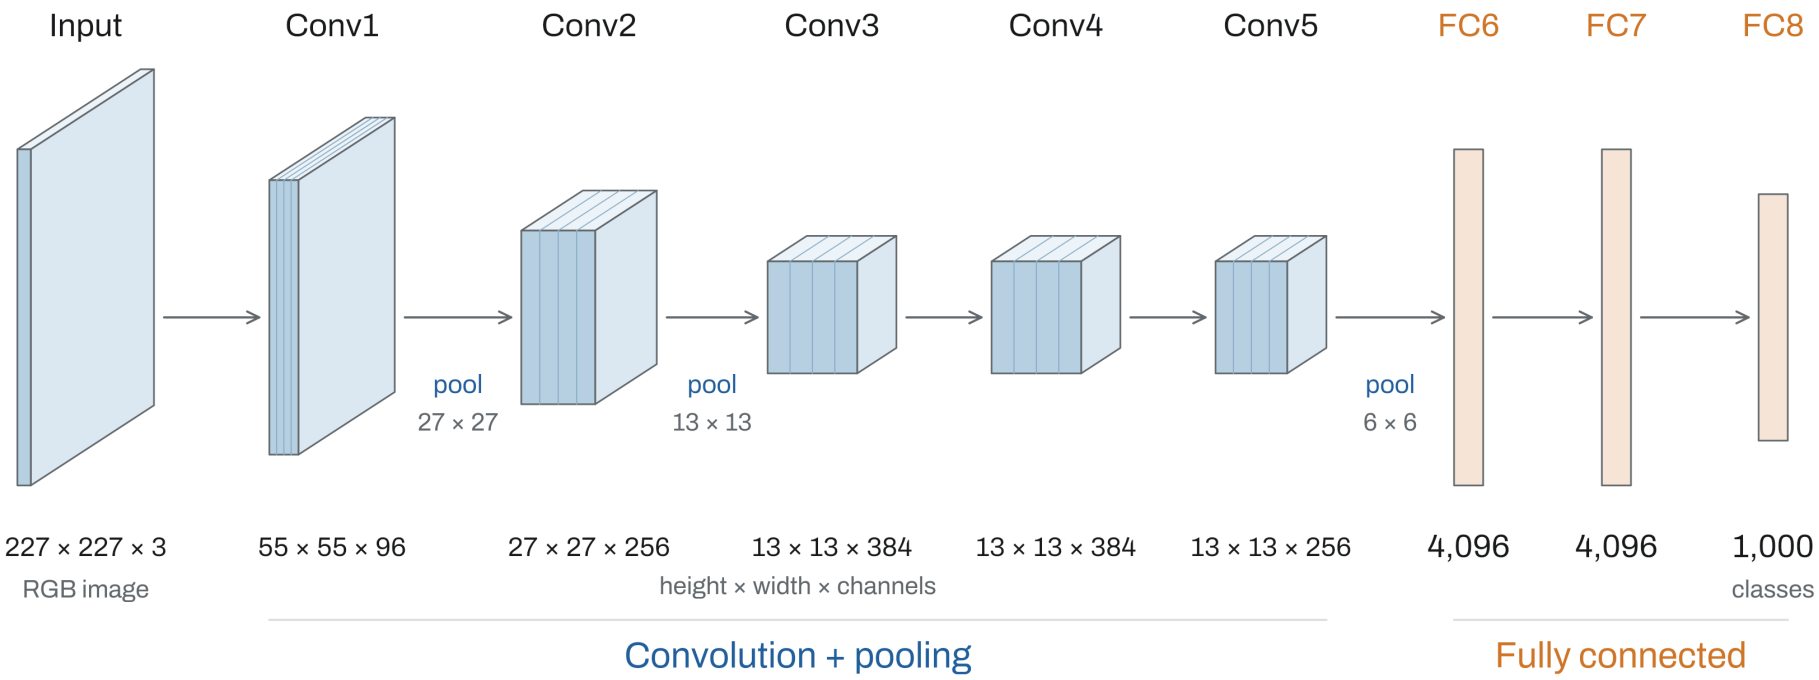

ImageNet 2012 winner: 15.3% top-5 error

Krizhevsky, Sutskever & Hinton (2012) · ImageNet classification with deep CNNs

Notes

Compare a visual model with cortex

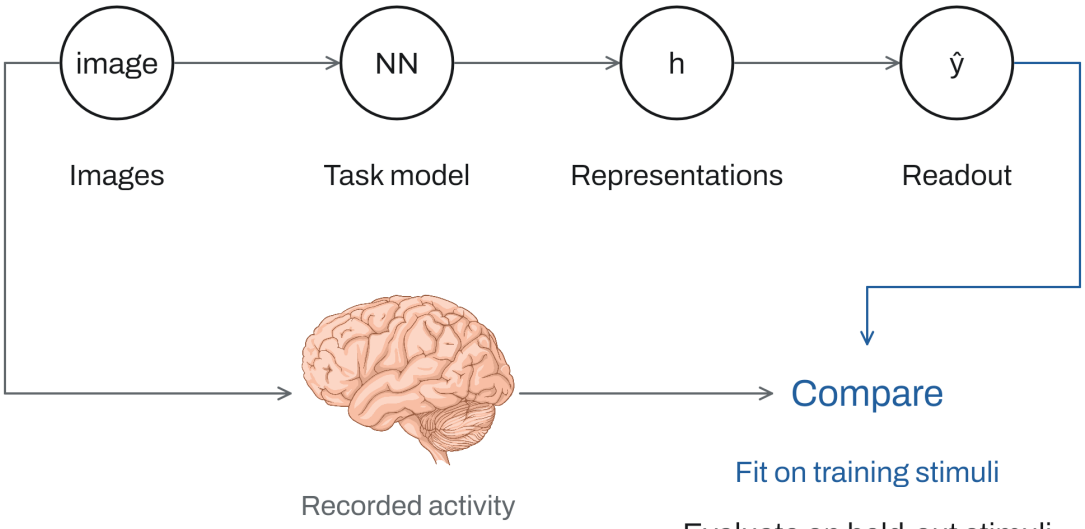

Fit on training stimuli

Evaluate on held-out stimuli

Optimize the model for object recognition.

Fit a readout to predict the recorded neural responses.

Compare predictions with recordings for held-out images.

Why test the readout on different images?

To test generalization instead of measuring how well it fits the training images.

Yamins et al. (2014) · PNAS · doi:10.1073/pnas.1403112111

Notes

Task performance and IT predictivity

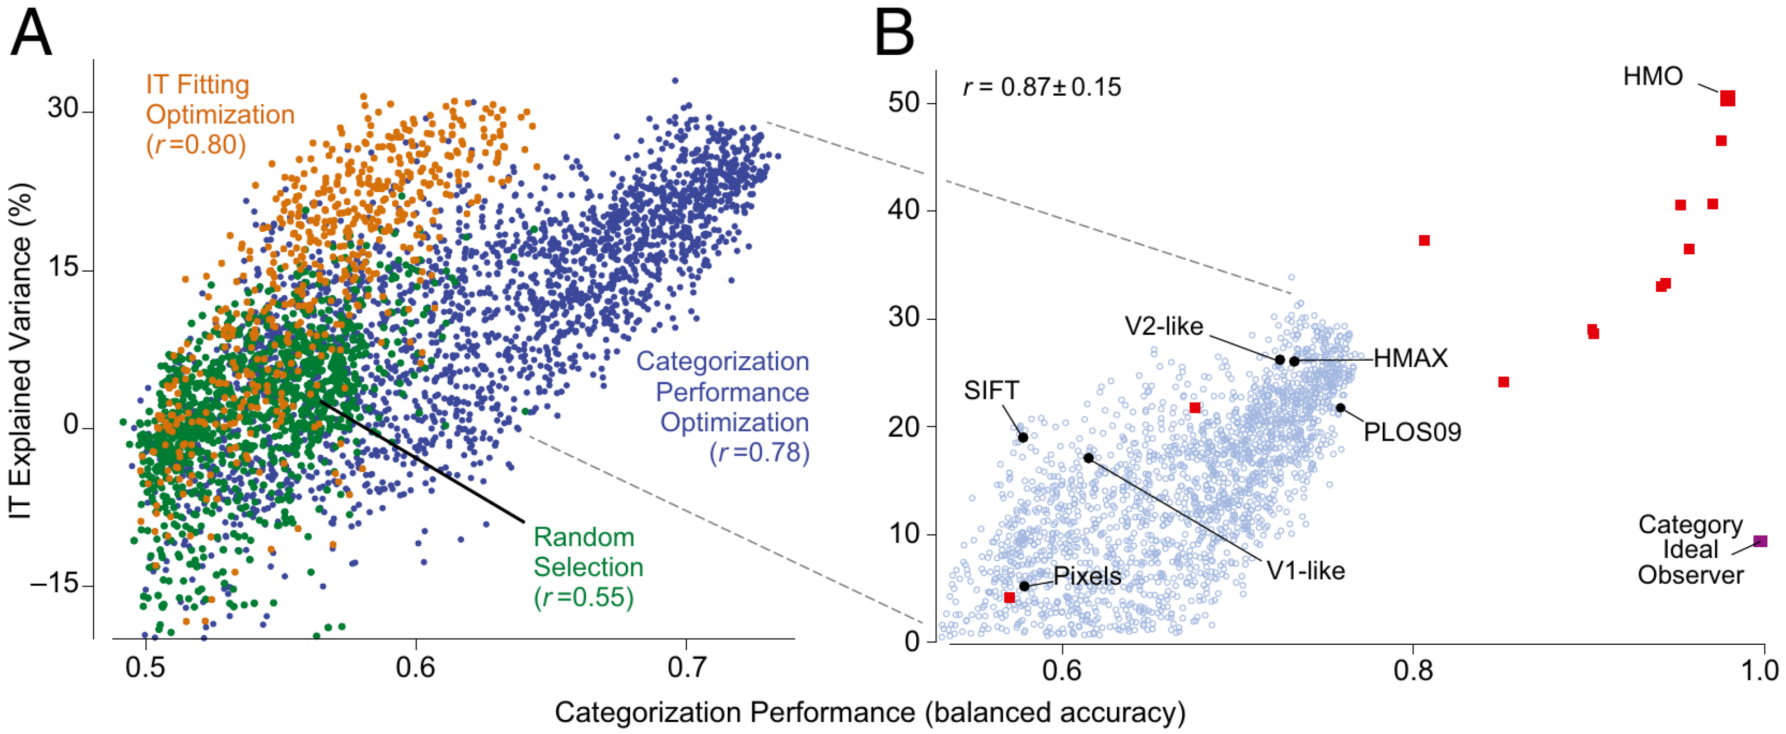

Yamins et al. (2014) · PNAS · doi:10.1073/pnas.1403112111

Notes

What the visual-cortex result supports

In these models, stronger object recognition tended to accompany better IT predictions.

Task optimization can help identify useful representations.

A neural fit does not establish identical biological mechanisms.

Yamins et al. (2014) · PNAS · doi:10.1073/pnas.1403112111

Notes

Transformers: represent a word using its context

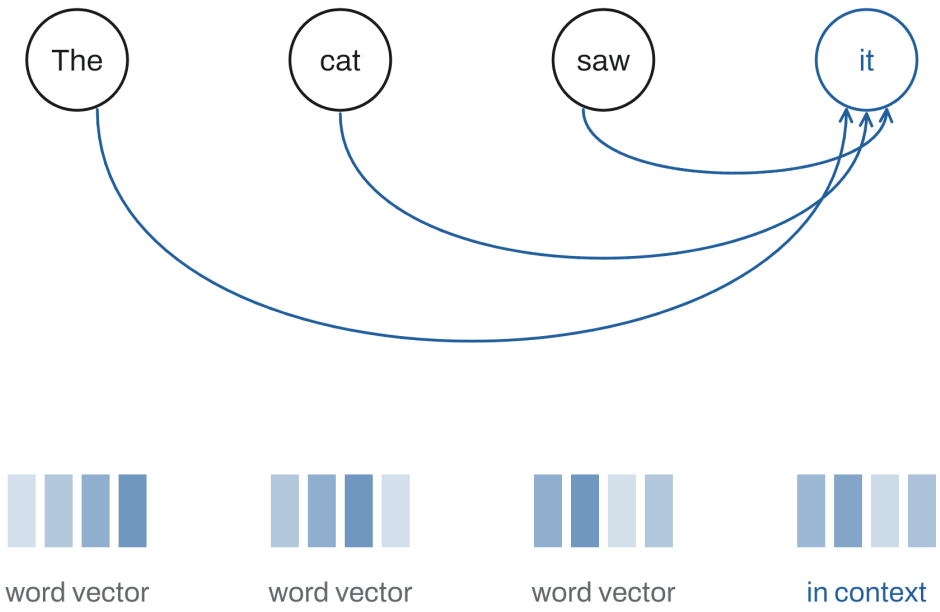

Start with a vector for each word.

Attention combines information from the available words.

Notes

Words become vectors; positions matter

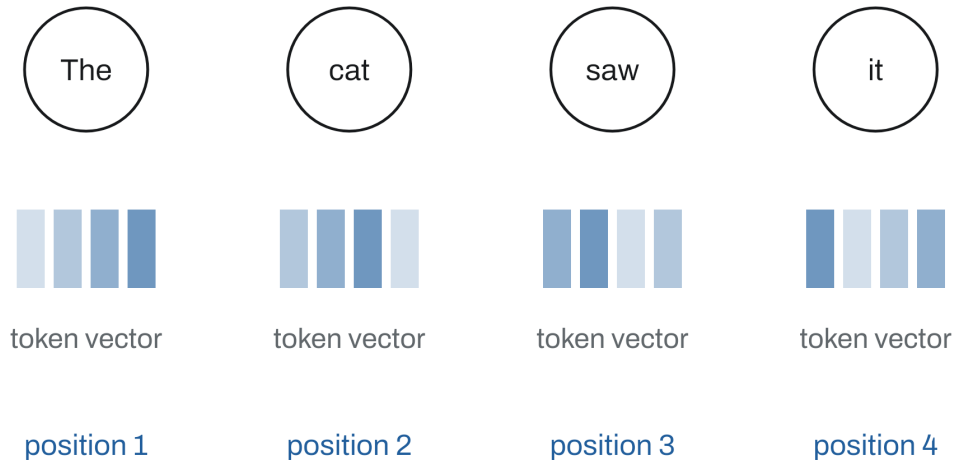

A token starts with a learned vector.

Position information tells the model about order.

Together they enter the first transformer block.

Radford et al. (2019) · GPT-2

Notes

A token asks for useful context

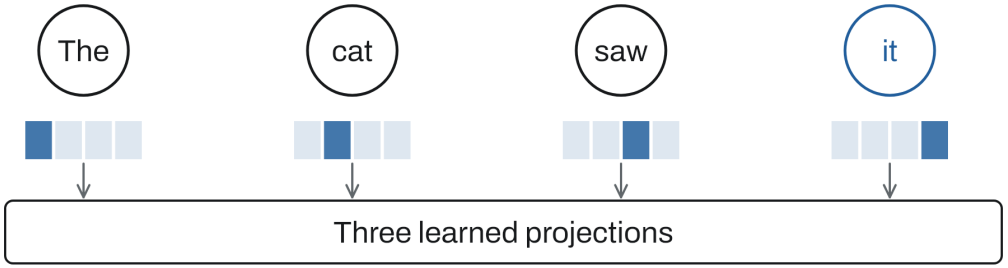

	The	cat	saw	it
Q	[0, 0]	[1, 0]	[0, 1]	[1, 0]
K	[0, 0]	[2, 0]	[0, 1]	[1, 0]
V	[0, 0]	[1, 0]	[0, 1]	[0.5, 0.5]

Query: what this token looks for.

Key: what each available token offers for matching.

Value: the information each token contributes.

Query token: it

Hand-set projections

Notes

Turn match scores into a weighted mixture

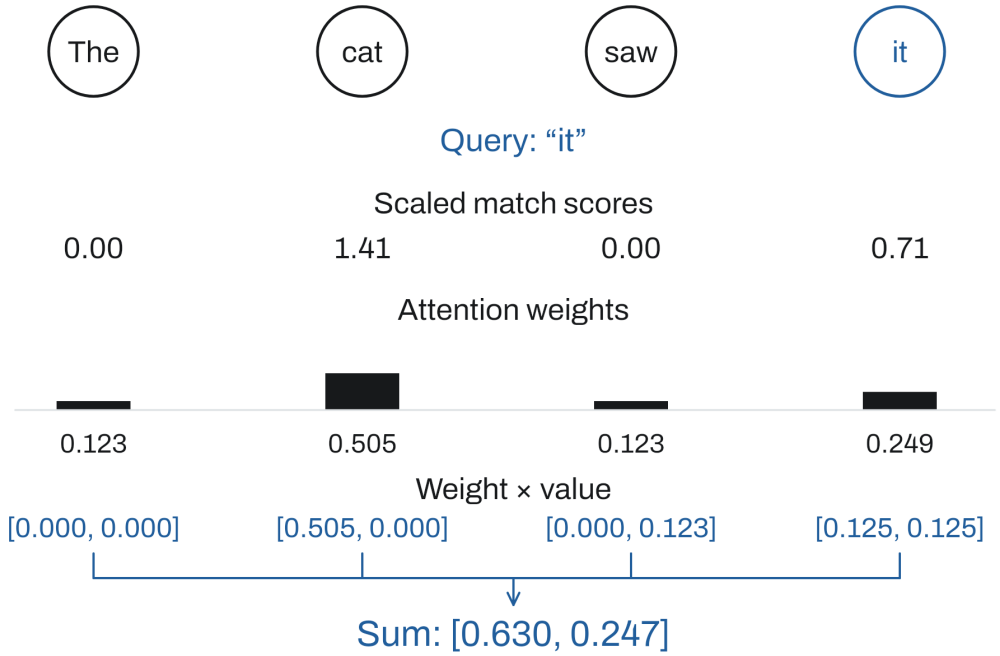

Compare the query with available keys.

Softmax turns scores into weights that sum to 1.

Use those weights to mix the value vectors.

Query token: it

Hand-set projections

Does the output simply copy the value for “cat”?

No. It combines the available values according to their weights.

Notes

Next-word prediction: no peeking ahead

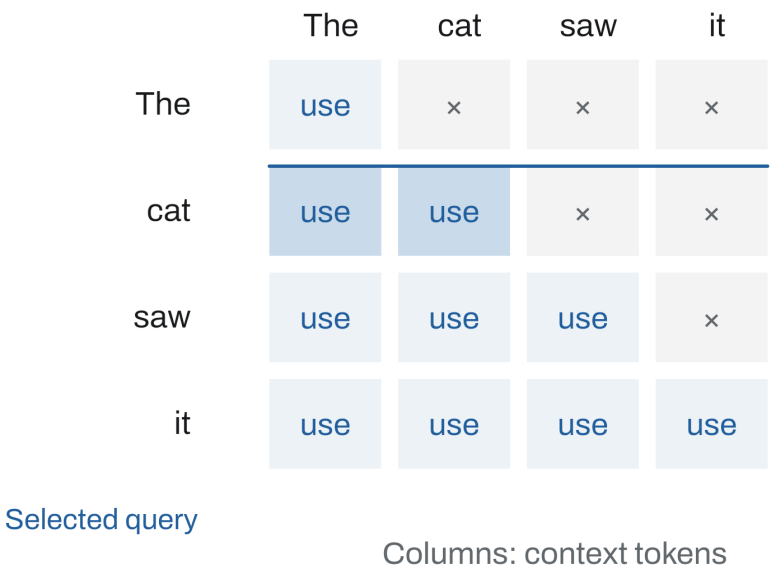

Causal attention uses only the current and earlier words.

Future positions are excluded before softmax.

Query row: cat

When predicting after “The cat”, may the model inspect “saw”?

No. That future token must be hidden for this prediction.

Notes

Attention is one part of a block

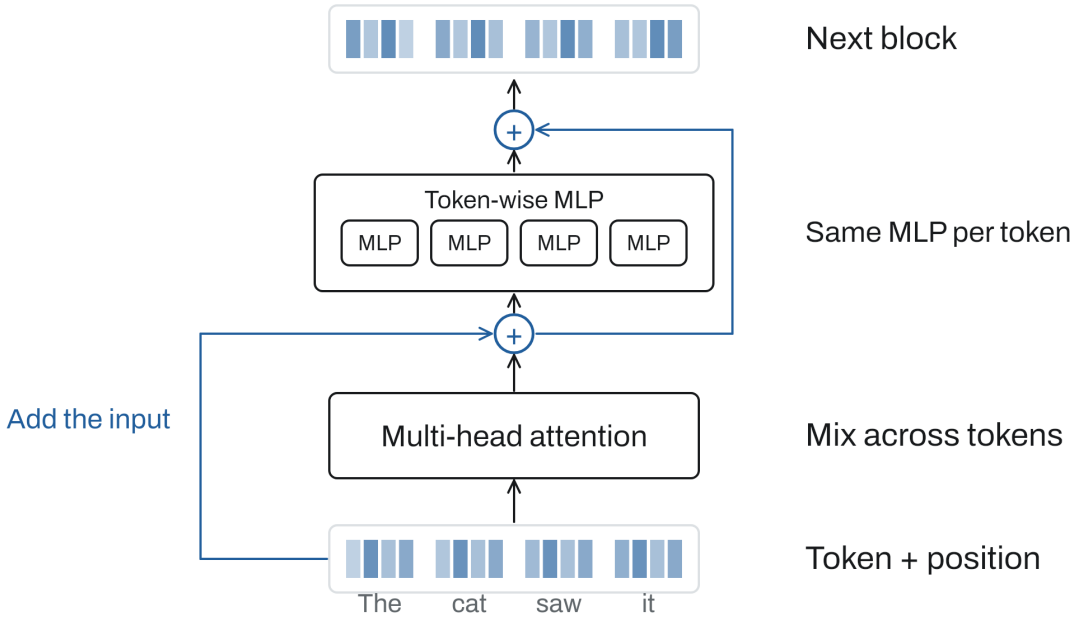

Attention mixes information across available positions.

A feedforward network transforms each position.

A residual path keeps the previous representation and adds an update.

Vaswani et al. (2017); Radford et al. (2019)

Notes

Several heads, several learned mixtures

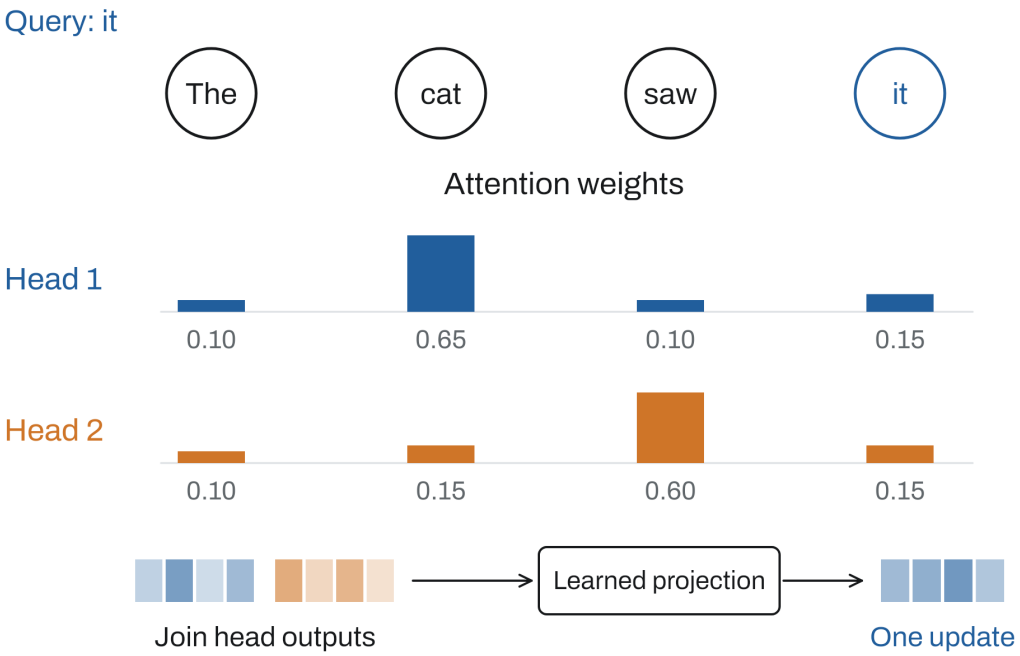

Each head has its own query, key and value projections.

Their outputs are combined.

Head roles are learned rather than assigned in advance.

Notes

Transformer demos

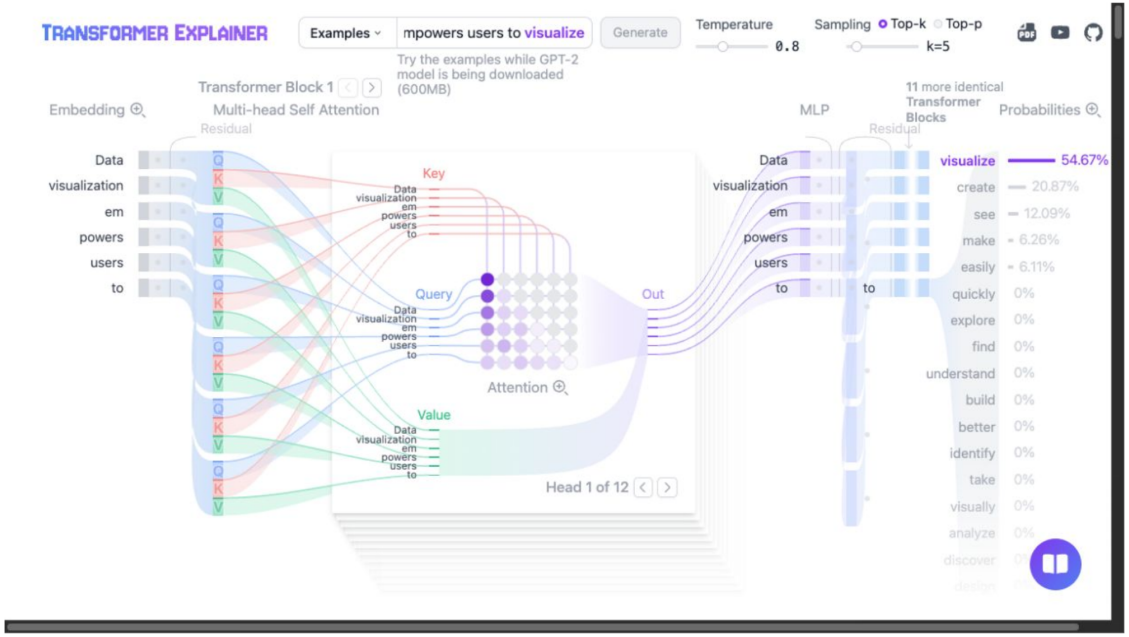

Transformer Explainer

Interactive class demo · GPT-2

3Blue1Brown

Attention, step by step · video

BertViz

Explore attention heads

Transformer Explainer · Georgia Tech Polo Club

Notes

Compare language models with brain responses

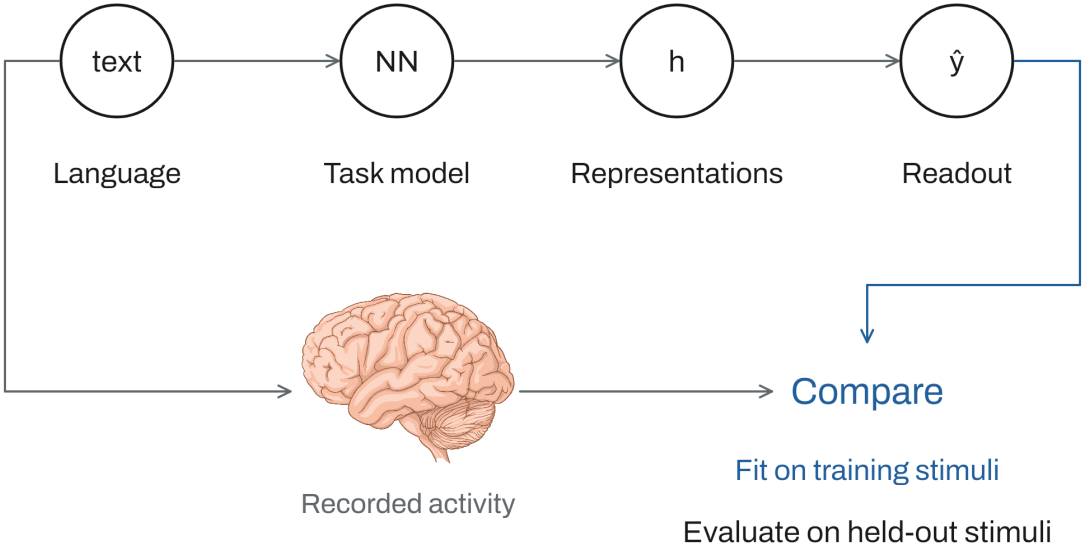

Present the same language to model and participants.

Fit a readout from representations to predicted responses.

Compare predictions with recordings for held-out stimuli.

Schrimpf et al. (2021) · PNAS · doi:10.1073/pnas.2105646118

Notes

Next-word performance and brain predictivity

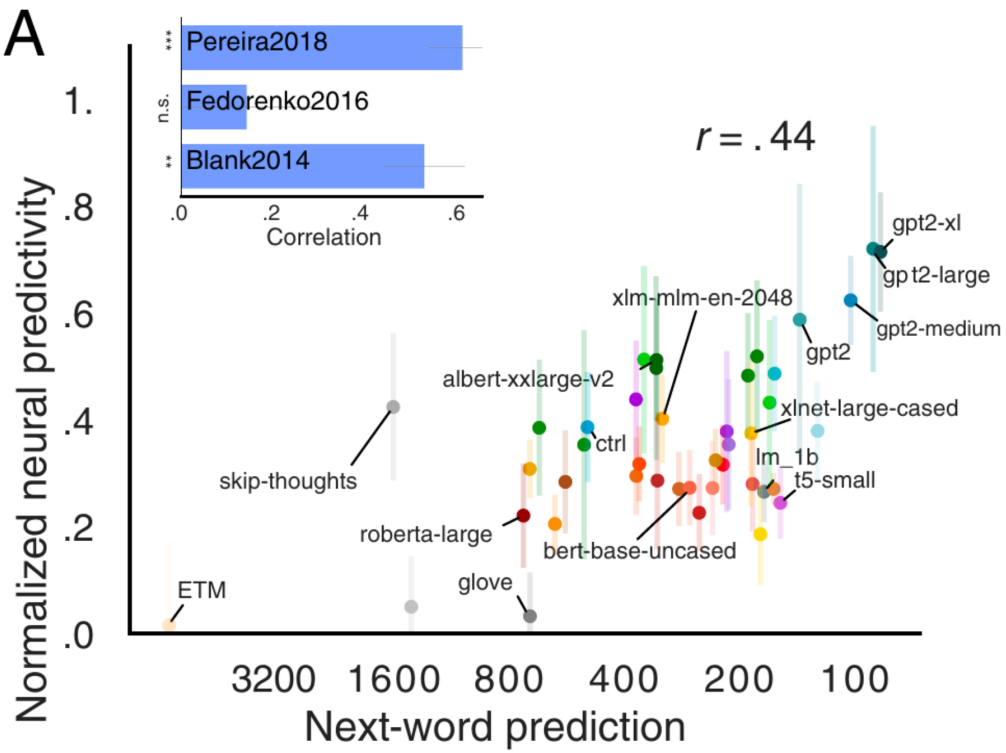

Perplexity scores next-word prediction: lower is better. It decreases from left to right.

Higher on the vertical axis means better predictions of recorded brain responses.

Each point represents a model.

Schrimpf et al. (2021) · PNAS · doi:10.1073/pnas.2105646118

Notes

A predictive model is a testable hypothesis

Better next-word prediction was associated with better neural fits.

Architecture and training both mattered.

A good fit does not prove the brain implements GPT.

Schrimpf et al. (2021) · PNAS · doi:10.1073/pnas.2105646118

Notes

Choose the information flow

Layers

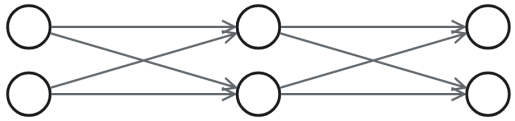

MLP: combine input features.

Recurrence

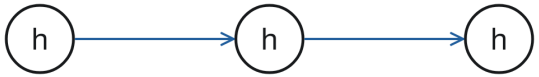

RNN / LSTM / GRU: update a state through time.

Convolution

CNN: reuse local filters across space.

Attention

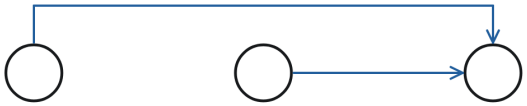

Transformer: mix available context with attention.

Notes

Neural networks as model organisms

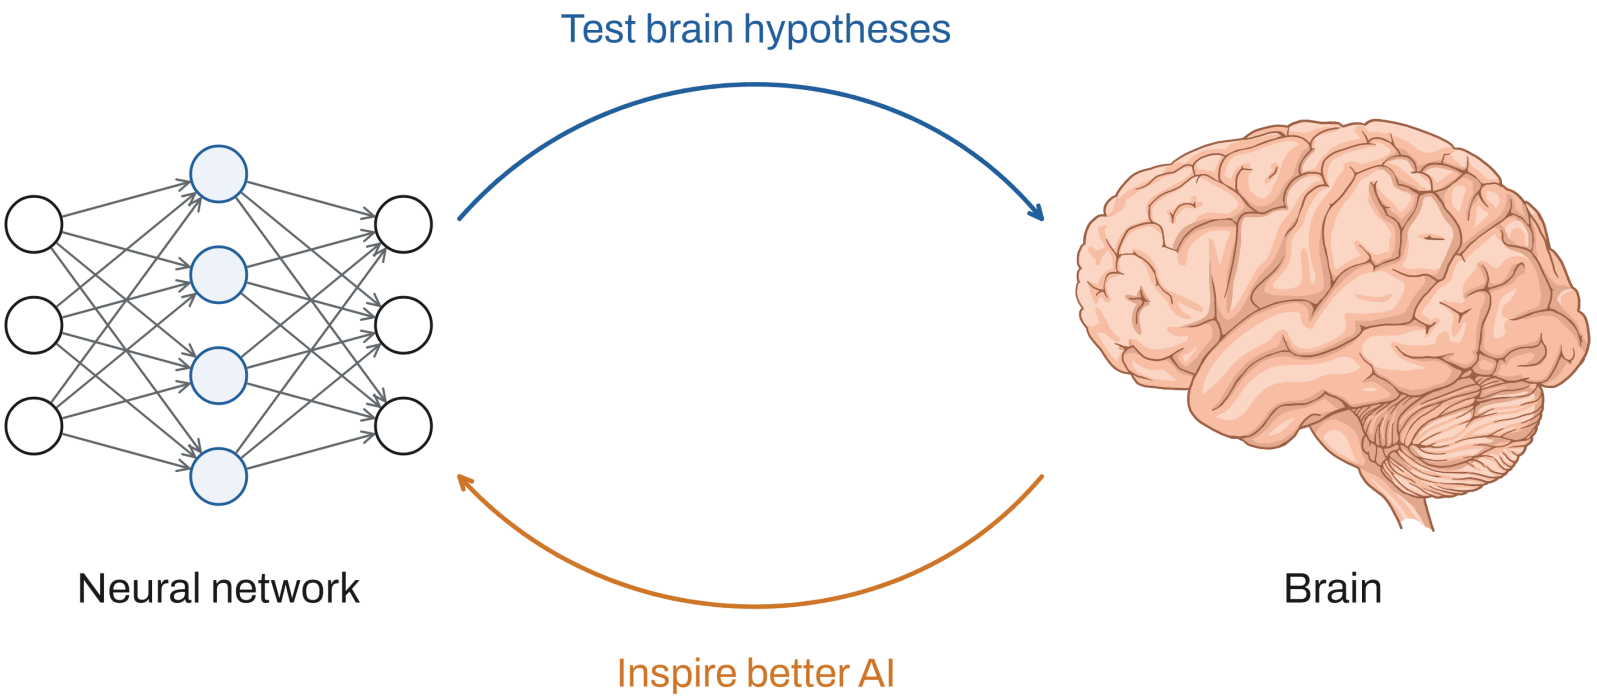

Notes

References

Christopher Olah (2015) · Understanding LSTM Networks

Rogers & McClelland (2004) · Semantic Cognition

Saffran, Aslin & Newport (1996) · Statistical learning in infants

Elman (1990) · Finding Structure in Time

Howard & Kahana (2002) · Temporal Context Model

Hochreiter & Schmidhuber (1997) · LSTM

Gers, Schmidhuber & Cummins (2000) · Learning to Forget

Cho et al. (2014) · RNN Encoder–Decoder / GRU

Yamins et al. (2014) · Visual cortex

Vaswani et al. (2017) · Attention Is All You Need

Radford et al. (2019) · GPT-2

Schrimpf et al. (2021) · Language and brain

Devlin et al. (2019) · BERT

Tulving (1972/1987) · Episodic and semantic memory

Mandler & McDonough (1993) · Concept Formation in Infancy

Rogers et al. (2004) · Semantic dementia

Saxe, McClelland & Ganguli (2019) · Semantic development

Anderson (2026) · Beni greenlet news report

Krizhevsky, Sutskever & Hinton (2012) · AlexNet

Richards et al. (2019) · Deep learning for neuroscience

Hassabis et al. (2017) · Neuroscience-inspired AI

NIAID Visual & Medical Arts (2024) · Human brain illustration · Courtesy of NIAID

Notes